CVE Lakehouse Assignment – Final Report

1. Introduction
This report documents the end-to-end implementation of a CVE (Common Vulnerabilities and Exposures) Lakehouse analytics pipeline using Databricks Community Edition. The project transforms raw JSON vulnerability records from the CVE cvelistV5 repository into structured Bronze and Silver Delta tables, followed by Gold-level SQL analytics generating insights about vendors, severity distribution, and vulnerability disclosure timelines.

2. Bronze Layer – Raw Ingestion
The Bronze layer loads raw JSON CVE records after downloading and extracting the cvelistV5 repository. In Databricks CE, ingestion was performed using */databricks/driver* storage to avoid DBFS restrictions. After extraction, recursive JSON loading produced a raw DataFrame, filtered to records published in 2024. Data quality checks validated record volume, null IDs, and uniqueness.

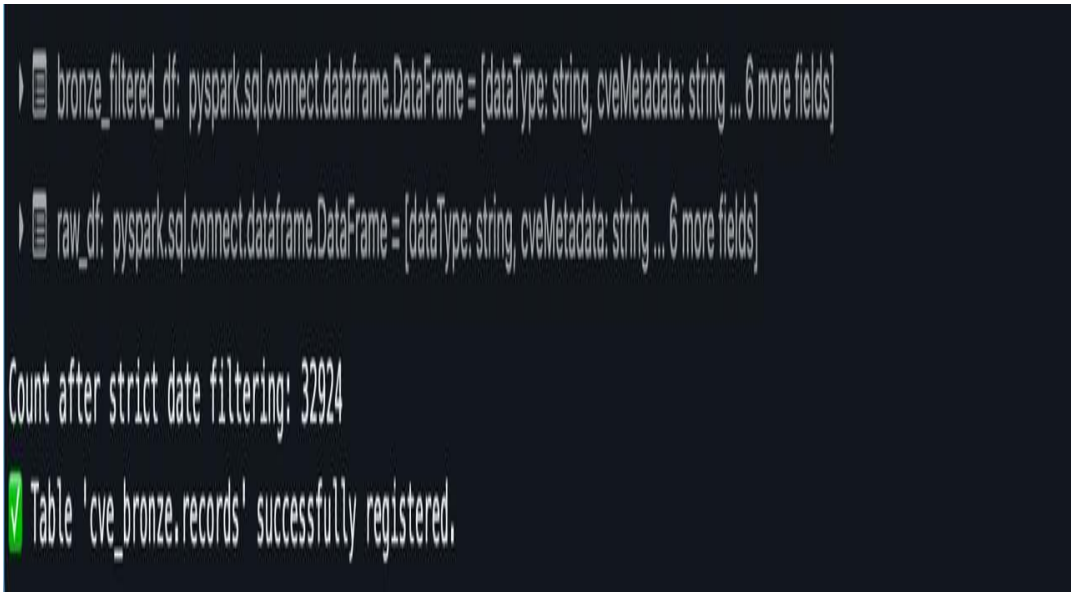

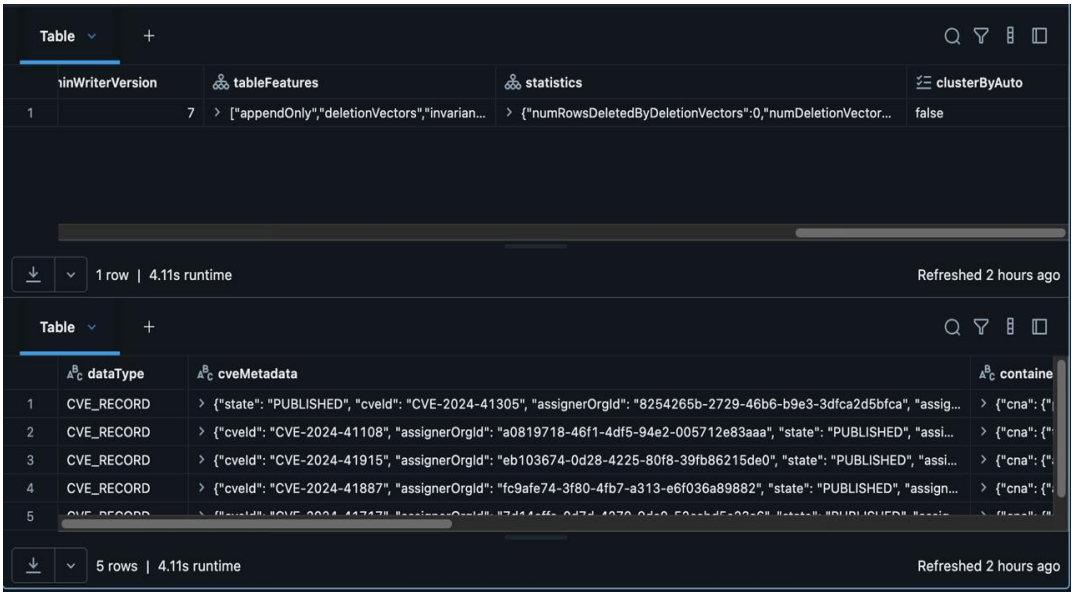

3. Silver Layer – Normalized Tables
The Silver layer transforms the Bronze dataset into two analytical tables:
• **core** – Flattened metadata including CVE ID, publication dates, CVSS score, severity, and description.
• **affected** – Exploded vendor–product associations for each CVE.
These normalized tables enable vendor-level analytics and severity scoring.

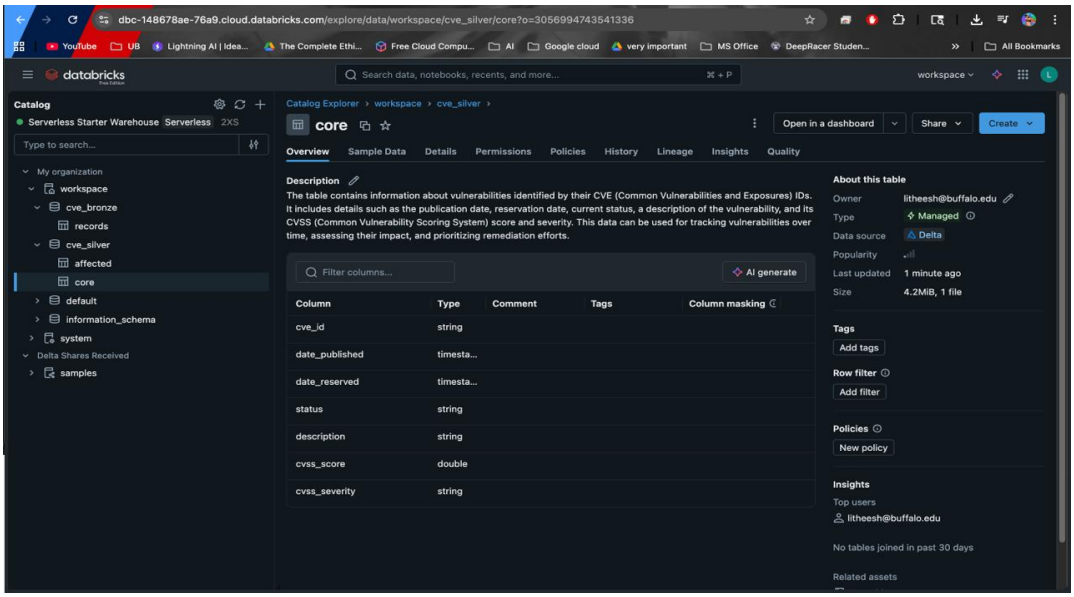

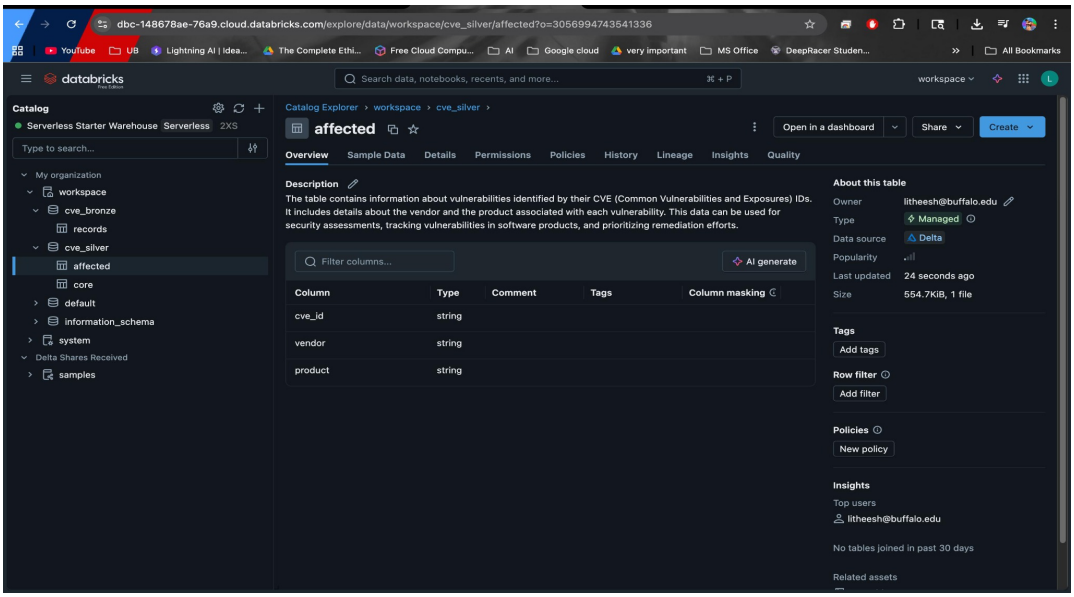

4. Gold Layer – SQL Analytics (EDA)

SQL was applied on Silver tables to generate insights including:

• Top affected vendors

• Severity distribution

• Disclosure lag analysis

These insights help identify high-risk vendors, vulnerability severity proportions, and the efficiency of disclosure.

Disclosure Lag Analysis (Top 5 Slowest Disclosures):

Table

	cve_id	date_reserved	date_published	disclosure_lag_days
1	CVE-2024-21635	2023-12-29T03:00:44.956+00:00	2025-11-14T14:11:38.230+00:00	686
2	CVE-2024-0028	2023-11-16T22:58:45.676+00:00	2025-09-05T16:10:01.094+00:00	659
3	CVE-2024-25621	2024-02-08T22:26:33.511+00:00	2025-11-06T18:36:21.566+00:00	637
4	CVE-2024-21927	2024-01-03T16:43:09.233+00:00	2025-09-23T21:33:54.121+00:00	629
5	CVE-2024-21935	2024-01-03T16:43:14.976+00:00	2025-09-23T21:38:22.057+00:00	629

5 rows | 10.79s runtime

Refreshed 2 hours ago

Summary Statistics for Disclosure Lag:

Table

	summary	disclosure_lag_days
1	count	38320
2	mean	50.82562630480167
3	stddev	75.92830749799869
4	min	0
5	max	686

5 rows | 10.79s runtime

Refreshed 2 hours ago

Key Insight: The average disclosure lag for 2024 CVEs is: 50.83 days.

CVSS Severity Distribution (Risk Bucketing):

Table

	cvss_severity	count	percentage
1	null	16555	42.71927334657962
2	MEDIUM	11795	30.436353314582096
3	HIGH	7588	19.58041958041958
4	CRITICAL	1788	4.613836348153692
5	LOW	1015	2.6191520656465306
6	NONE	12	0.03096534461848115

6 rows | 1.92s runtime

Refreshed 2 hours ago

Key Insight: There are 1,788 CRITICAL severity vulnerabilities in the dataset.

Top 10 Affected Vendors by Total CVE Count:

Table +

	vendor	total_cves
1	Microsoft	13161
2	n/a	6591
3	Linux	6152
4	Brother Industries, Ltd	4427
5	Red Hat	3913
6	Siemens	2545
7	Apple	1692
8	Unknown	1092
9	Lenovo	929
10	Adobe	751

10 rows | 1.53s runtime

Refreshed 2 hours ago

🔑 Key Insight: The top vendor affected by the highest number of 2024 CVEs is Microsoft with 13,161 total affected products/versions.

5. Conclusion

This assignment demonstrates a complete, production-inspired Lakehouse pipeline using Databricks CE. Despite storage restrictions, the workflow successfully implemented:

• Bronze ingestion using driver storage
• Silver normalization (core + affected)
• Gold-level analytics producing actionable cybersecurity insights

The structured pipeline is scalable, reproducible, and aligns with modern Delta Lake best practices.